

DS project Group-12

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.feature_selection import f_regression, SelectKBest
from sklearn.ensemble import ExtraTreesRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import RandomizedSearchCV
```

```
In [2]: df = pd.read_csv('Crime_Incidents_in_2012.csv')
# df2 = pd.read_csv('Crime_Incidents_in_2013.csv')
# df3 = pd.read_csv('Crime_Incidents_in_2014.csv')
# df4 = pd.read_csv('Crime_Incidents_in_2015.csv')
# df5 = pd.read_csv('Crime_Incidents_in_2016.csv')
# df6 = pd.read_csv('Crime_Incidents_in_2017.csv')
# df7 = pd.read_csv('Crime_Incidents_in_2018.csv')
# df8 = pd.read_csv('Crime_Incidents_in_2019.csv')
# df9 = pd.read_csv('Crime_Incidents_in_2020.csv')
# df0 = pd.read_csv('Crime_Incidents_in_2021.csv')
```

```
In [3]: # COLUMN_NAMES = ['X', 'Y', 'CCN', 'REPORT_DAT', 'SHIFT', 'METHOD', 'OFFENSE', 'BLOCK',
#                        'YBLOCK', 'WARD', 'ANC', 'DISTRICT', 'PSA', 'NEIGHBORHOOD_CLUSTER',
#                        'VOTING_PRECINCT', 'LATITUDE', 'LONGITUDE', 'BID', 'START_DATE', 'END_DATE']
# df = pd.DataFrame(columns=COLUMN_NAMES)
# df = pd.concat([df1, df2, df3, df4, df5, df6, df7, df8, df9, df0], ignore_index=True)
```

```
In [4]: # df.to_csv('data.csv', index=False)
```

```
In [5]: #df=pd.read_csv('data.csv')
```

```
In [6]: df=df.iloc[:500]
df.shape
```

```
Out[6]: (500, 24)
```

```
In [7]: df.columns
```


```
Out[7]: Index(['X', 'Y', 'CCN', 'REPORT_DAT', 'SHIFT', 'METHOD', 'OFFENSE', 'BLOCK',
              'XBLOCK', 'YBLOCK', 'WARD', 'ANC', 'DISTRICT', 'PSA',
              'NEIGHBORHOOD_CLUSTER', 'BLOCK_GROUP', 'CENSUS_TRACT',
              'VOTING_PRECINCT', 'LATITUDE', 'LONGITUDE', 'BID', 'START_DATE',
              'END_DATE', 'OBJECTID'],
              dtype='object')
```

```
In [8]: df.head()
```

Out[8]:

| | X | Y | CCN | REPORT_DAT | SHIFT | METHOD | OFFENSE | BLOCK | XBL |
|---|------------|-----------|----------|--------------------------|----------|--------|-----------|---|-----|
| 0 | -76.999518 | 38.901924 | 9074624 | 2012-04-25T00:00:00.000Z | MIDNIGHT | OTHERS | SEX ABUSE | 900 - 999 BLOCK OF 5TH STREET NE | 40 |
| 1 | -76.995541 | 38.905032 | 10123633 | 2012-02-29T00:00:00.000Z | MIDNIGHT | OTHERS | SEX ABUSE | 700 - 799 BLOCK OF FLORIDA AVENUE NE | 40 |
| 2 | -76.948897 | 38.885680 | 11102619 | 2012-05-14T00:00:00.000Z | MIDNIGHT | GUN | HOMICIDE | 153 - 399 BLOCK OF RIDGE ROAD SE | 40 |
| 3 | -76.967571 | 38.855724 | 11141272 | 2012-06-25T00:00:00.000Z | MIDNIGHT | OTHERS | HOMICIDE | 2800 - 2899 BLOCK OF BUENA VISTA TERRACE SE | 40 |
| 4 | -76.939620 | 38.910718 | 11158196 | 2012-01-05T00:00:00.000Z | MIDNIGHT | OTHERS | HOMICIDE | 4280 - 4499 BLOCK OF DOUGLAS STREET NE | 40 |

5 rows × 24 columns



```
In [9]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   X                                      500 non-null    float64
1   Y                                      500 non-null    float64
2   CCN                                    500 non-null    int64
3   REPORT_DAT                            500 non-null    object
4   SHIFT                                 500 non-null    object
5   METHOD                                 500 non-null    object
6   OFFENSE                               500 non-null    object
7   BLOCK                                 500 non-null    object
8   XBLOCK                                500 non-null    int64
9   YBLOCK                                500 non-null    int64
10  WARD                                   500 non-null    float64
11  ANC                                    500 non-null    object
12  DISTRICT                              500 non-null    float64
13  PSA                                    500 non-null    float64
14  NEIGHBORHOOD_CLUSTER                 492 non-null    object
15  BLOCK_GROUP                          499 non-null    object
16  CENSUS_TRACT                         499 non-null    float64
17  VOTING_PRECINCT                      500 non-null    object
18  LATITUDE                             500 non-null    float64
19  LONGITUDE                             500 non-null    float64
20  BID                                   76 non-null     object
21  START_DATE                           500 non-null    object
22  END_DATE                              496 non-null    object
23  OBJECTID                             500 non-null    int64
dtypes: float64(8), int64(4), object(12)
memory usage: 93.9+ KB

```

In [10]: `df.describe()`

Out[10]:

| | X | Y | CCN | XBLOCK | YBLOCK | WARD | DISTRICT | |
|--------------|------------|------------|--------------|---------------|---------------|------------|------------|-----|
| count | 500.000000 | 500.000000 | 5.000000e+02 | 500.000000 | 500.000000 | 500.000000 | 500.000000 | 500 |
| mean | -77.007607 | 38.905599 | 1.199100e+07 | 399341.046000 | 137526.588000 | 4.634000 | 3.80600 | 380 |
| std | 0.036525 | 0.034332 | 1.732524e+05 | 3167.849485 | 3811.243612 | 2.310567 | 1.99407 | 190 |
| min | -77.091050 | 38.813478 | 9.074624e+06 | 392105.000000 | 127300.000000 | 1.000000 | 1.00000 | 100 |
| 25% | -77.032420 | 38.887592 | 1.200530e+07 | 397190.750000 | 135527.750000 | 2.000000 | 2.00000 | 200 |
| 50% | -77.008114 | 38.905650 | 1.200648e+07 | 399296.000000 | 137532.000000 | 5.000000 | 4.00000 | 400 |
| 75% | -76.983647 | 38.927541 | 1.200982e+07 | 401419.750000 | 139962.250000 | 7.000000 | 5.25000 | 600 |
| max | -76.917789 | 38.986543 | 1.201348e+07 | 407132.000000 | 146512.000000 | 8.000000 | 7.00000 | 700 |

In [11]: `df.isnull().sum()`

```
Out[11]: X 0
Y 0
CCN 0
REPORT_DAT 0
SHIFT 0
METHOD 0
OFFENSE 0
BLOCK 0
XBLOCK 0
YBLOCK 0
WARD 0
ANC 0
DISTRICT 0
PSA 0
NEIGHBORHOOD_CLUSTER 8
BLOCK_GROUP 1
CENSUS_TRACT 1
VOTING_PRECINCT 0
LATITUDE 0
LONGITUDE 0
BID 424
START_DATE 0
END_DATE 4
OBJECTID 0
dtype: int64
```

```
In [12]: #df.drop(['BID', 'OCTO_RECORD_ID'], axis=1, inplace=True)
df.drop(['BID'], axis=1, inplace=True)
# having more than 75% null values
```

```
In [13]: df.shape
```

```
Out[13]: (500, 23)
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: X 0
Y 0
CCN 0
REPORT_DAT 0
SHIFT 0
METHOD 0
OFFENSE 0
BLOCK 0
XBLOCK 0
YBLOCK 0
WARD 0
ANC 0
DISTRICT 0
PSA 0
NEIGHBORHOOD_CLUSTER 8
BLOCK_GROUP 1
CENSUS_TRACT 1
VOTING_PRECINCT 0
LATITUDE 0
LONGITUDE 0
START_DATE 0
END_DATE 4
OBJECTID 0
dtype: int64
```

```
In [15]: df.dropna(subset=['WARD', 'DISTRICT', 'PSA', 'NEIGHBORHOOD_CLUSTER', 'BLOCK_GROUP', 'CENSUS'])
```

```
In [16]: df.isnull().sum()
```

```
Out[16]: X 0
Y 0
CCN 0
REPORT_DAT 0
SHIFT 0
METHOD 0
OFFENSE 0
BLOCK 0
XBLOCK 0
YBLOCK 0
WARD 0
ANC 0
DISTRICT 0
PSA 0
NEIGHBORHOOD_CLUSTER 0
BLOCK_GROUP 0
CENSUS_TRACT 0
VOTING_PRECINCT 0
LATITUDE 0
LONGITUDE 0
START_DATE 0
END_DATE 0
OBJECTID 0
dtype: int64
```

```
In [17]: df.shape
```

```
Out[17]: (487, 23)
```

```
In [18]: df['VOTING_PRECINCT'].head()
```

```
Out[18]: 0    Precinct 83
          1    Precinct 83
          3    Precinct 134
          4    Precinct 92
          5    Precinct 5
          Name: VOTING_PRECINCT, dtype: object
```

```
In [19]: df['NEIGHBORHOOD_CLUSTER'].head()
```

```
Out[19]: 0    Cluster 25
          1    Cluster 25
          3    Cluster 36
          4    Cluster 29
          5    Cluster 4
          Name: NEIGHBORHOOD_CLUSTER, dtype: object
```

Removing "Precinct" and "Cluster" from each record of VOTING_PRECINCT and NEIGHBORHOOD_CLUSTER..

```
In [20]: df['VOTING_PRECINCT']=df['VOTING_PRECINCT'].str.replace('Precinct ','').astype(int)
```

```
In [21]: df['NEIGHBORHOOD_CLUSTER']=df['NEIGHBORHOOD_CLUSTER'].str.replace('Cluster ','').astype(int)
```

```
In [22]: df['VOTING_PRECINCT'].head()
```

```
Out[22]: 0      83
          1      83
          3     134
          4      92
          5       5
          Name: VOTING_PRECINCT, dtype: int32
```

```
In [23]: df['NEIGHBORHOOD_CLUSTER'].head()
```

```
Out[23]: 0      25
          1      25
          3      36
          4      29
          5       4
          Name: NEIGHBORHOOD_CLUSTER, dtype: int32
```

Values in column X and Y is same as that of the column LATITUDE and LONGITUDE. So we can remove X and Y.

```
In [24]: df.drop(['X','Y'],axis=1,inplace=True)
```

```
In [25]: df.head()
```

| Out[25]: | CCN | REPORT_DAT | SHIFT | METHOD | OFFENSE | BLOCK | XBLOCK | YBLOCK | WAF |
|----------|----------|--------------------------|----------|--------|-------------|---|--------|--------|-----|
| 0 | 9074624 | 2012-04-25T00:00:00.000Z | MIDNIGHT | OTHERS | SEX ABUSE | 900 - 999 BLOCK OF 5TH STREET NE | 400042 | 137118 | € |
| 1 | 10123633 | 2012-02-29T00:00:00.000Z | MIDNIGHT | OTHERS | SEX ABUSE | 700 - 799 BLOCK OF FLORIDA AVENUE NE | 400387 | 137463 | € |
| 3 | 11141272 | 2012-06-25T00:00:00.000Z | MIDNIGHT | OTHERS | HOMICIDE | 2800 - 2899 BLOCK OF BUENA VISTA TERRACE SE | 402815 | 131990 | € |
| 4 | 11158196 | 2012-01-05T00:00:00.000Z | MIDNIGHT | OTHERS | HOMICIDE | 4280 - 4499 BLOCK OF DOUGLAS STREET NE | 405237 | 138096 | 7 |
| 5 | 12005414 | 2012-01-11T18:52:00.000Z | EVENING | OTHERS | THEFT/OTHER | 1500 - 1599 BLOCK OF 32ND STREET NW | 394480 | 138000 | 2 |

5 rows × 21 columns

Also dates are represented as object type, it should be converted into datetime datatype.

```

In [26]: df['START_DATE']=pd.to_datetime(df['START_DATE'])

In [27]: df['END_DATE']=pd.to_datetime(df['END_DATE'])

In [28]: df['REPORT_DAT']=pd.to_datetime(df['REPORT_DAT'])

In [29]: df['OFFENSE'].unique()

Out[29]: array(['SEX ABUSE', 'HOMICIDE', 'THEFT/OTHER',
        'ASSAULT W/DANGEROUS WEAPON', 'ROBBERY', 'THEFT F/AUTO',
        'MOTOR VEHICLE THEFT', 'BURGLARY'], dtype=object)

In [30]: # Import Label encoder
from sklearn import preprocessing as pp

```

```
offense_encoder = pp.LabelEncoder()
df['OFFENSE']=offense_encoder.fit_transform(df['OFFENSE'])
```

```
In [31]: df['OFFENSE'].head()
```

```
Out[31]: 0    5
1    5
3    2
4    2
5    7
Name: OFFENSE, dtype: int32
```

Same like that convert BLOCK,ANC column to int32

```
In [32]: df.head()
```

```
Out[32]:
```

| | CCN | REPORT_DAT | SHIFT | METHOD | OFFENSE | BLOCK | XBLOCK | YBLOCK | WARD | AI |
|---|----------|------------------------------|----------|--------|---------|---|--------|--------|------|----|
| 0 | 9074624 | 2012-04-25 00:00:00+00:00 | MIDNIGHT | OTHERS | 5 | 900 - 999 BLOCK OF 5TH STREET NE | 400042 | 137118 | 6.0 | |
| 1 | 10123633 | 2012-02-29 00:00:00+00:00 | MIDNIGHT | OTHERS | 5 | 700 - 799 BLOCK OF FLORIDA AVENUE NE | 400387 | 137463 | 6.0 | |
| 3 | 11141272 | 2012-06-25 00:00:00+00:00 | MIDNIGHT | OTHERS | 2 | 2800 - 2899 BLOCK OF BUENA VISTA TERRACE SE | 402815 | 131990 | 8.0 | |
| 4 | 11158196 | 2012-01-05 00:00:00+00:00 | MIDNIGHT | OTHERS | 2 | 4280 - 4499 BLOCK OF DOUGLAS STREET NE | 405237 | 138096 | 7.0 | |
| 5 | 12005414 | 2012-01-11 18:52:00+00:00 | EVENING | OTHERS | 7 | 1500 - 1599 BLOCK OF 32ND STREET NW | 394480 | 138000 | 2.0 | |

5 rows × 21 columns

```
In [33]: df['BLOCK'].unique().size
```


Out[33]: 459

```
In [34]: block_encoder = pp.LabelEncoder()
df['BLOCK']=block_encoder.fit_transform(df['BLOCK'])

anc_encoder = pp.LabelEncoder()
df['ANC']=anc_encoder.fit_transform(df['ANC'])
```

```
In [35]: df['BLOCK'].head()
```

```
Out[35]: 0    423
1    397
3    218
4    316
5    102
Name: BLOCK, dtype: int32
```

```
In [36]: df['ANC'].head()
```

```
Out[36]: 0     27
1     27
3     36
4     32
5      8
Name: ANC, dtype: int32
```

```
In [37]: df.head()
```

```
Out[37]:
```

| | CCN | REPORT_DAT | SHIFT | METHOD | OFFENSE | BLOCK | XBLOCK | YBLOCK | WARD | ANC |
|---|----------|------------------------------|----------|--------|---------|-------|--------|--------|------|-----|
| 0 | 9074624 | 2012-04-25 00:00:00+00:00 | MIDNIGHT | OTHERS | 5 | 423 | 400042 | 137118 | 6.0 | 27 |
| 1 | 10123633 | 2012-02-29 00:00:00+00:00 | MIDNIGHT | OTHERS | 5 | 397 | 400387 | 137463 | 6.0 | 27 |
| 3 | 11141272 | 2012-06-25 00:00:00+00:00 | MIDNIGHT | OTHERS | 2 | 218 | 402815 | 131990 | 8.0 | 36 |
| 4 | 11158196 | 2012-01-05 00:00:00+00:00 | MIDNIGHT | OTHERS | 2 | 316 | 405237 | 138096 | 7.0 | 32 |
| 5 | 12005414 | 2012-01-11 18:52:00+00:00 | EVENING | OTHERS | 7 | 102 | 394480 | 138000 | 2.0 | 8 |

5 rows × 21 columns

One Hot Encoding SHIFT and METHOD using get_dummies() function. It will create new columns with value 0 and 1 only.

```
In [38]: df['SHIFT'].unique()
```

```
Out[38]: array(['MIDNIGHT', 'EVENING', 'DAY'], dtype=object)
```

```
In [39]: df['METHOD'].unique()
```

```
Out[39]: array(['OTHERS', 'GUN', 'KNIFE'], dtype=object)
```

```
In [40]: shift_encoder = pp.LabelEncoder()
df['SHIFT']=shift_encoder.fit_transform(df['SHIFT'])

met_encoder = pp.LabelEncoder()
df['METHOD']=met_encoder.fit_transform(df['METHOD'])
```

```
In [41]: df.head()
```

```
Out[41]:
```

| | CCN | REPORT_DAT | SHIFT | METHOD | OFFENSE | BLOCK | XBLOCK | YBLOCK | WARD | ANC | ... |
|---|----------|------------------------------|-------|--------|---------|-------|--------|--------|------|-----|-----|
| 0 | 9074624 | 2012-04-25 00:00:00+00:00 | 2 | 2 | 5 | 423 | 400042 | 137118 | 6.0 | 27 | ... |
| 1 | 10123633 | 2012-02-29 00:00:00+00:00 | 2 | 2 | 5 | 397 | 400387 | 137463 | 6.0 | 27 | ... |
| 3 | 11141272 | 2012-06-25 00:00:00+00:00 | 2 | 2 | 2 | 218 | 402815 | 131990 | 8.0 | 36 | ... |
| 4 | 11158196 | 2012-01-05 00:00:00+00:00 | 2 | 2 | 2 | 316 | 405237 | 138096 | 7.0 | 32 | ... |
| 5 | 12005414 | 2012-01-11 18:52:00+00:00 | 1 | 2 | 7 | 102 | 394480 | 138000 | 2.0 | 8 | ... |

5 rows × 21 columns

```
In [42]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 487 entries, 0 to 499
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   CCN                                  487 non-null    int64
1   REPORT_DAT                          487 non-null    datetime64[ns, UTC]
2   SHIFT                              487 non-null    int32
3   METHOD                              487 non-null    int32
4   OFFENSE                             487 non-null    int32
5   BLOCK                              487 non-null    int32
6   XBLOCK                             487 non-null    int64
7   YBLOCK                             487 non-null    int64
8   WARD                               487 non-null    float64
9   ANC                                487 non-null    int32
10  DISTRICT                           487 non-null    float64
11  PSA                                487 non-null    float64
12  NEIGHBORHOOD_CLUSTER               487 non-null    int32
13  BLOCK_GROUP                       487 non-null    object
14  CENSUS_TRACT                      487 non-null    float64
15  VOTING_PRECINCT                   487 non-null    int32
16  LATITUDE                          487 non-null    float64
17  LONGITUDE                         487 non-null    float64
18  START_DATE                        487 non-null    datetime64[ns, UTC]
19  END_DATE                          487 non-null    datetime64[ns, UTC]
20  OBJECTID                          487 non-null    int64
dtypes: datetime64[ns, UTC](3), float64(6), int32(7), int64(4), object(1)
memory usage: 70.4+ KB

```

Now here block group is only object left , should convert it to int.

```
In [43]: df['BLOCK_GROUP'].unique().size
```

```
Out[43]: 250
```

```
In [44]: bg_encoder = pp.LabelEncoder()
df['BLOCK_GROUP'] = bg_encoder.fit_transform(df['BLOCK_GROUP'])
```

```
In [45]: df['BLOCK_GROUP'].head()
```

```
Out[45]: 0    242
1    241
3    139
4    207
5         2
Name: BLOCK_GROUP, dtype: int32
```

```
In [46]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 487 entries, 0 to 499
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CCN                                    487 non-null    int64
1   REPORT_DAT                            487 non-null    datetime64[ns, UTC]
2   SHIFT                                 487 non-null    int32
3   METHOD                                 487 non-null    int32
4   OFFENSE                               487 non-null    int32
5   BLOCK                                 487 non-null    int32
6   XBLOCK                                487 non-null    int64
7   YBLOCK                                487 non-null    int64
8   WARD                                  487 non-null    float64
9   ANC                                   487 non-null    int32
10  DISTRICT                              487 non-null    float64
11  PSA                                   487 non-null    float64
12  NEIGHBORHOOD_CLUSTER                 487 non-null    int32
13  BLOCK_GROUP                          487 non-null    int32
14  CENSUS_TRACT                         487 non-null    float64
15  VOTING_PRECINCT                      487 non-null    int32
16  LATITUDE                             487 non-null    float64
17  LONGITUDE                             487 non-null    float64
18  START_DATE                           487 non-null    datetime64[ns, UTC]
19  END_DATE                             487 non-null    datetime64[ns, UTC]
20  OBJECTID                             487 non-null    int64
dtypes: datetime64[ns, UTC](3), float64(6), int32(8), int64(4)
memory usage: 68.5 KB

```

Successfully all the features converted to int,float,datetime

Now should remove unwanted features like CCN, strat_date, end_date, OBJECTID. Because we dont want CCN, OBJECTID and also instead of two dates we have another feature report_date.

```
In [47]: df.drop(['CCN', 'START_DATE', 'END_DATE', 'OBJECTID'], axis=1, inplace=True)
```

Report date converted to month, day day of week.

```
In [48]: df['MONTH'] = df['REPORT_DAT'].dt.month
df['DAY'] = df['REPORT_DAT'].dt.day
df['Day_Week'] = df['REPORT_DAT'].dt.dayofweek
df.drop('REPORT_DAT', axis=1, inplace=True)
```

Final DataSet

```
In [49]: df.head()
```

```
Out[49]:
```

| | SHIFT | METHOD | OFFENSE | BLOCK | XBLOCK | YBLOCK | WARD | ANC | DISTRICT | PSA | NEIGHBORI |
|---|-------|--------|---------|-------|--------|--------|------|-----|----------|-------|-----------|
| 0 | 2 | 2 | 5 | 423 | 400042 | 137118 | 6.0 | 27 | 1.0 | 104.0 | |
| 1 | 2 | 2 | 5 | 397 | 400387 | 137463 | 6.0 | 27 | 5.0 | 506.0 | |
| 3 | 2 | 2 | 2 | 218 | 402815 | 131990 | 8.0 | 36 | 7.0 | 702.0 | |
| 4 | 2 | 2 | 2 | 316 | 405237 | 138096 | 7.0 | 32 | 6.0 | 601.0 | |
| 5 | 1 | 2 | 7 | 102 | 394480 | 138000 | 2.0 | 8 | 2.0 | 206.0 | |

```
In [50]: df.shape
```

```
Out[50]: (487, 19)
```

```
In [51]: df.info()
```

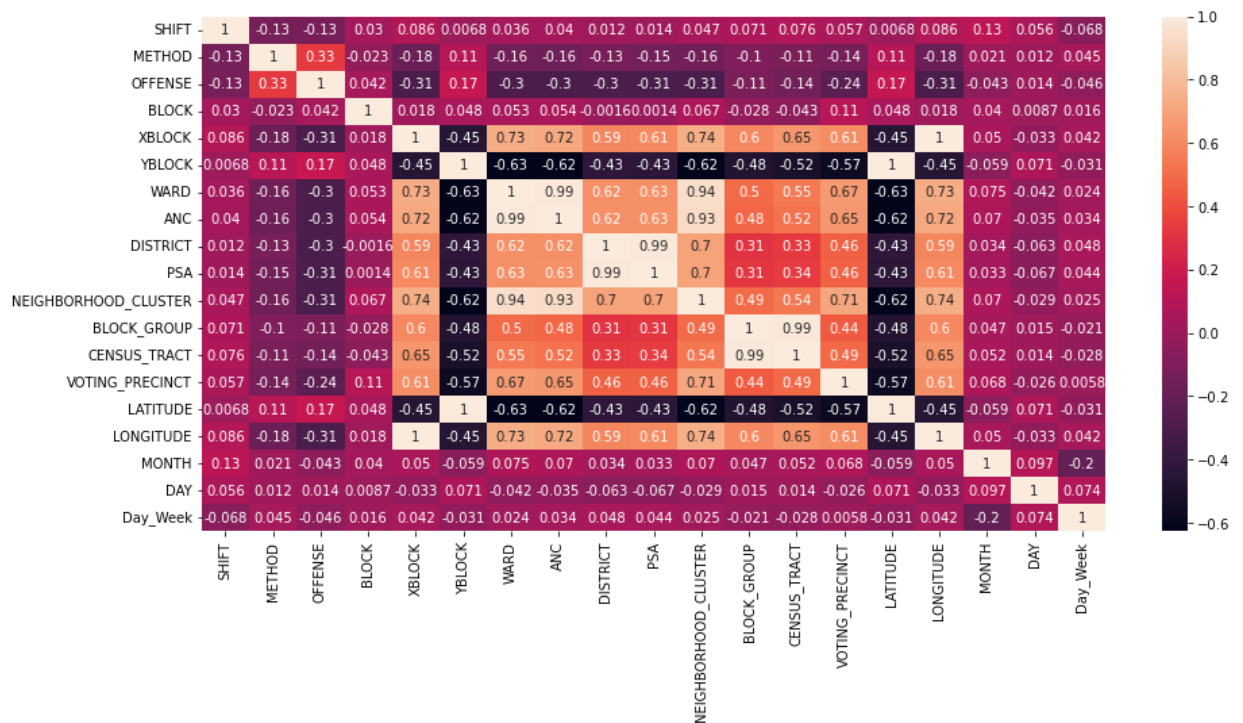
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 487 entries, 0 to 499
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   SHIFT                 487 non-null    int32
1   METHOD                 487 non-null    int32
2   OFFENSE               487 non-null    int32
3   BLOCK                 487 non-null    int32
4   XBLOCK                487 non-null    int64
5   YBLOCK                487 non-null    int64
6   WARD                  487 non-null    float64
7   ANC                   487 non-null    int32
8   DISTRICT              487 non-null    float64
9   PSA                   487 non-null    float64
10  NEIGHBORHOOD_CLUSTER  487 non-null    int32
11  BLOCK_GROUP           487 non-null    int32
12  CENSUS_TRACT          487 non-null    float64
13  VOTING_PRECINCT       487 non-null    int32
14  LATITUDE              487 non-null    float64
15  LONGITUDE             487 non-null    float64
16  MONTH                 487 non-null    int64
17  DAY                   487 non-null    int64
18  Day_Week              487 non-null    int64
dtypes: float64(6), int32(8), int64(5)
memory usage: 60.9 KB
```

```
In [52]: df.iloc[0]
```

```
Out[52]: SHIFT                2.000000
METHOD                2.000000
OFFENSE              5.000000
BLOCK                423.000000
XBLOCK              400042.000000
YBLOCK              137118.000000
WARD                 6.000000
ANC                 27.000000
DISTRICT             1.000000
PSA                 104.000000
NEIGHBORHOOD_CLUSTER 25.000000
BLOCK_GROUP          242.000000
CENSUS_TRACT        10600.000000
VOTING_PRECINCT      83.000000
LATITUDE             38.901916
LONGITUDE            -76.999516
MONTH                4.000000
DAY                 25.000000
Day_Week             2.000000
Name: 0, dtype: float64
```

```
In [53]: # correlation heat map
plt.figure(figsize=[15,7])
sb.heatmap(df.corr(), annot=True)
```


```
Out[53]: <AxesSubplot:>
```



```
In [54]: # Important feature using ExtraTreesRegressor
from sklearn.ensemble import ExtraTreesRegressor
X = df.drop('OFFENSE',axis=1)
y = df['OFFENSE']
model = ExtraTreesRegressor()
model.fit(X,y)
```

```
Out[54]: ExtraTreesRegressor()
```

```
In [55]: # plot graph of feature importances for 6 better visualization
plt.figure(figsize=[12,6])
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(6).plot(kind='barh')
plt.show()
```



```
In [56]: print(feat_importances.sort_values(ascending=False))
```

```
METHOD          0.144868
DAY              0.079951
BLOCK            0.077631
Day_Week         0.071847
XBLOCK           0.068694
LONGITUDE        0.066604
ANC              0.057553
PSA              0.054855
LATITUDE         0.049692
YBLOCK           0.048032
SHIFT            0.045858
BLOCK_GROUP      0.043801
NEIGHBORHOOD_CLUSTER 0.042991
VOTING_PRECINCT  0.041342
DISTRICT         0.039411
CENSUS_TRACT     0.037104
WARD             0.028830
MONTH            0.000935
dtype: float64
```

```
In [57]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2, random_state=365)
```

```
In [58]: from sklearn.preprocessing import LabelEncoder,StandardScaler
from sklearn.model_selection import train_test_split,cross_val_score,KFold
#Importing Models
from sklearn.linear_model import LogisticRegression,LinearRegression,Lasso,Ridge,Bayes
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn import svm
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.metrics import r2_score, mean_squared_error
import seaborn as sns
```

```
In [59]: models = {
    'Logistic' : LogisticRegression(),
    'random forest' : RandomForestRegressor(),
    'decision tree' : DecisionTreeRegressor(max_depth=5),
    'support vector': svm.SVC(),
    'knn' : KNeighborsRegressor(n_neighbors = 4)
}
```

```
In [60]: for name, model in models.items():
    model.fit(X_train, y_train)
    print(f'{name} trained')
```

C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

`n_iter_i = _check_optimize_result(`

Logistic trained

random forest trained

decision tree trained

support vector trained

knn trained

```
In [61]: results = {}
kf = KFold(n_splits= 10)

for name, model in models.items():
    result = np.mean(np.sqrt(-cross_val_score(model, X_train, y_train, scoring='neg_mean_squared_error', cv=kf)))
    results[name] = result
```



```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(
 C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
 _logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(
 C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
 _logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(
 C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
 _logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(
 C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model
 _logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```
In [62]: for name, result in results.items():
        print(f"{name} : {round(result, 3)}")
```

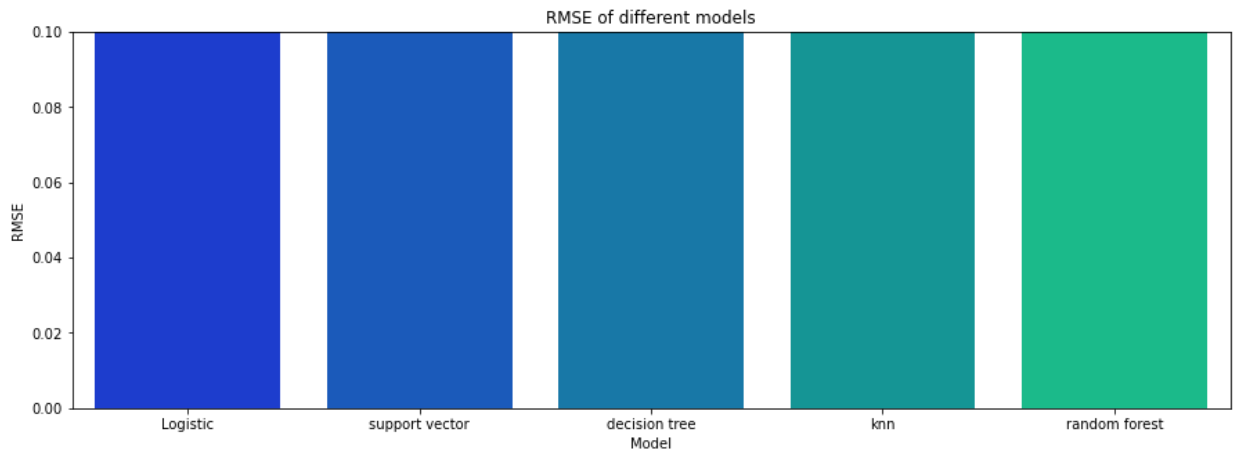
```
Logistic : 2.91
random forest : 2.125
decision tree : 2.348
support vector : 2.761
knn : 2.322
```

```
In [63]: results_df = pd.DataFrame(results, index=range(0,1)).T.rename(columns={0: 'RMSE'}).sort_index()
        results_df.T
```

```
Out[63]:
```

| | Logistic | support vector | decision tree | knn | random forest |
|------|----------|----------------|---------------|----------|---------------|
| RMSE | 2.909916 | 2.760594 | 2.347988 | 2.322145 | 2.125403 |

```
In [64]: plt.figure(figsize = (15, 5))
        sns.barplot(x= results_df.index, y = results_df['RMSE'], palette = 'winter')
        plt.ylim(0,0.1)
        plt.xlabel('Model')
        plt.ylabel('RMSE')
        plt.title('RMSE of different models');
```



```
In [65]: def prediction(model,X_train,y_train,X_test,y_test):
          model.fit(X_train,y_train)
          pred_data=np.exp(model.predict(X_test))
          return r2_score(np.exp(y_test),pred_data)
```

```
In [66]: for name,model in models.items():
          score=prediction(model,X_train,y_train,X_test,y_test)
          print(f'{name} r2_score is {score}')
```

C:\Users\91938\Documents\sample_project_1\env\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Logistic r2_score is -0.37982653533356814


random forest r2_score is -0.28386472502650695

decision tree r2_score is -0.22290812828469675


support vector r2_score is -0.042798579783665414

knn r2_score is -0.5942471882184519


```
In [67]: model=LinearRegression()
          model.fit(X_train,y_train)
          pred_y=np.exp(model.predict(X_test))
          sns.histplot(np.exp(y_test)-pred_y)
          plt.xlim(-0.1,0.1)
          plt.ylim(0,80)
          plt.show()
```




```
In [68]: model=RandomForestRegressor()  
model.fit(X_train,y_train)  
pred_y=np.exp(model.predict(X_test))  
sns.histplot(np.exp(y_test)-pred_y)  
plt.xlim(-0.1,0.1)  
plt.ylim(0,80)  
plt.show()
```




```
In [69]: model=DecisionTreeRegressor(max_depth=5)  
model.fit(X_train,y_train)  
pred_y=np.exp(model.predict(X_test))  
sns.histplot(np.exp(y_test)-pred_y)  
plt.xlim(-0.1,0.1)  
plt.ylim(0,80)  
plt.show()
```




```
In [70]: model=SVR()
model.fit(X_train,y_train)
pred_y=np.exp(model.predict(X_test))
sns.histplot(np.exp(y_test)-pred_y)
plt.xlim(-0.1,0.1)
plt.ylim(0,80)
plt.show()
```



```
In [71]: model=KNeighborsRegressor(n_neighbors = 4)
model.fit(X_train,y_train)
pred_y=np.exp(model.predict(X_test))
sns.histplot(np.exp(y_test)-pred_y)
plt.xlim(-0.1,0.1)
plt.ylim(0,80)
plt.show()
```



```
In [72]: # •Similarly plot execution time for different values of k.
k = []
acc = []
time = []
for i in range(1, 20):
    kn = KNeighborsRegressor(n_neighbors=i)
    kn.fit(X_train, y_train)
    kn_pred = kn.predict(X_test)
    kn_acc = r2_score(y_test, kn_pred)
    k.append(i)
    acc.append(kn_acc)
    time.append(1/i)
# Plotting accuracy and execution time for different values of k
plt.plot(k, acc, label='Accuracy')
plt.plot(k, time, label='Time')
plt.legend()
plt.show()
```



```
In [73]: from sklearn.tree import DecisionTreeRegressor
for i in range(1, 10):
    dt = DecisionTreeRegressor(max_depth=i)
    dt.fit(X_train, y_train)
    dt_pred = dt.predict(X_test)
    dt_acc = r2_score(y_test, dt_pred)
    print(dt_acc)
```

```


0.05230631512283146
0.01714083312990855
0.004949400115422331
-0.1092329926322293
-0.2709842509575755
-0.2774021692346915
-0.4608248117722109
-0.5487967648453589
-0.6432599642018022

```

```

In [74]: # Make predictions on dataset. Plot accuracy and time for varying parameters
#decision tree model
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score
# Checking accuracy of each model for various values of max_depth
acc = []
time = []
k = []
for i in range(1,20):
    dt = DecisionTreeRegressor(max_depth=i)
    dt.fit(X_train, y_train)
    dt_pred = dt.predict(X_test)
    dt_acc = r2_score(y_test, dt_pred)
    acc.append(dt_acc)
    k.append(i)
    time.append(1/i)
# Plotting accuracy of each model for various values of max_depth
plt.plot(k, acc, label='Accuracy')
plt.plot(k, time, label='Time')
plt.legend()
plt.show()

```




```


In [75]: k = []
acc = []
time = []
for i in range(1, 20):
    kn = SVR()
    kn.fit(X_train, y_train)
    kn_pred = kn.predict(X_test)
    kn_acc = r2_score(y_test, kn_pred)
    k.append(i)
    acc.append(kn_acc)
    time.append(1/i)
# Plotting accuracy and execution time for different values of k

```

```
plt.plot(k, acc, label='Accuracy')
plt.plot(k, time, label='Time')
plt.legend()
plt.show()
```



```
In [76]: k = []
acc = []
time = []
for i in range(1, 20):
    kn = RandomForestRegressor()
    kn.fit(X_train, y_train)
    kn_pred = kn.predict(X_test)
    kn_acc = r2_score(y_test, kn_pred)
    k.append(i)
    acc.append(kn_acc)
    time.append(1/i)
# Plotting accuracy and execution time for different values of k
plt.plot(k, acc, label='Accuracy')
plt.plot(k, time, label='Time')
plt.legend()
plt.show()
```




```
In [77]: k = []
acc = []
time = []
for i in range(1, 20):
    kn = RandomForestRegressor()
```



```

kn.fit(X_train, y_train)
kn_pred = kn.predict(X_test)
kn_acc = r2_score(y_test, kn_pred)
k.append(i)
acc.append(kn_acc)
time.append(1/i)
# Plotting accuracy and execution time for different values of k
plt.plot(k, acc, label='Accuracy')
plt.plot(k, time, label='Time')
plt.legend()
plt.show()

```



```

In [81]: from sklearn.metrics import mean_absolute_error
xgb_model = RandomForestRegressor()
xgb_model.fit(X_train, y_train)
final_preds = xgb_model.predict(X_test)

model_results = pd.DataFrame([y_test.values, final_preds])
model_results = model_results.transpose()
model_results = model_results.rename(columns={0: 'Actual', 1: 'Predicted'})

print(model_results.describe(), '\n')

```

| | Actual | Predicted |
|-------|-----------|-----------|
| count | 98.000000 | 98.000000 |
| mean | 5.132653 | 4.844592 |
| std | 2.064115 | 1.383657 |
| min | 0.000000 | 0.610000 |
| 25% | 4.000000 | 4.172500 |
| 50% | 6.000000 | 4.960000 |
| 75% | 7.000000 | 5.850000 |
| max | 7.000000 | 6.890000 |

```

In [82]: !pip install pandoc

```

```

Requirement already satisfied: pandoc in c:\users\91938\documents\sample_project_1\env\lib\site-packages (2.2)
Requirement already satisfied: plumbum in c:\users\91938\documents\sample_project_1\env\lib\site-packages (from pandoc) (1.8.0)
Requirement already satisfied: ply in c:\users\91938\documents\sample_project_1\env\lib\site-packages (from pandoc) (3.11)
Requirement already satisfied: pywin32 in c:\users\91938\documents\sample_project_1\env\lib\site-packages (from plumbum->pandoc) (302)

```

